

Last Time:

- LQR as a QP
- Riccati Recursion

Today:

- Infinite-Horizon LQR
 - Controllability
 - Dynamic Programming
-

* Infinite Horizon LQR:

- For time-invariant LQR, K matrices converge to constant values
- For stabilization problems we usually use constant K
- Backward recursion for P :
$$K_n = (R + B^T P_{n+1} B)^{-1} B^T P_{n+1} A$$
$$P_n = Q + A^T P_{n+1} (A - B K_n)$$
- Infinite horizon limit $\Rightarrow P_n = P_{n+1} = P_{\text{inf}}$
 \Rightarrow Solve as a root-finding / fixed point problem for P_{inf}
- Julia/MATLAB dare function does this for you

* Controllability

- How do we know if LQR will work?
- We already know $Q \geq 0, R \geq 0$
- For the time-invariant case there is a simple answer.
- For any initial state x_0, x_N is given by:

$$x_N = Ax_{N-1} + Bu_{N-1}$$

$$= A(Ax_{N-2} + Bu_{N-2}) + Bu_{N-1}$$

⋮

$$= A^N x_0 + A^{N-1}Bu_0 + \dots + Bu_{N-1}$$

$$= \underbrace{\begin{bmatrix} B & AB & A^2B & \dots & A^{N-1}B \end{bmatrix}}_{\text{"Controllability Matrix"}} \begin{bmatrix} u_{N-1} \\ u_{N-2} \\ \vdots \\ u_0 \end{bmatrix} + A^N x_0$$

- This is equivalent to a linear least-squares problem for $u_{0:N-1}$:

$$\begin{bmatrix} u_{N-1} \\ \vdots \\ u_0 \end{bmatrix} = \underbrace{\left[C^\top \left(CC^\top \right)^{-1} \right]}_{\text{"Pseudo inverse"}} (x_N - A^N x_0)$$

must be invertable

"Pseudo inverse"

- For CCT to be invertible:

$$\Rightarrow \text{rank}(C) = n \quad (n = \dim(X))$$

- I get to stop at n time steps because the Cayley-Hamilton theorem says that A^n can be written as a linear combination of lower powers of A :

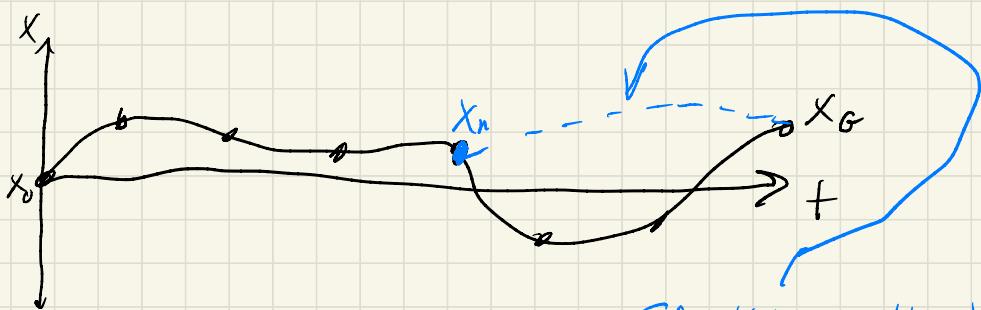
$$A^n = \sum_{k=0}^{n-1} \alpha_k A^k \quad (\text{for some } \alpha_k)$$

- Therefore adding more time steps / columns to C can't increase the rank

$$\Rightarrow C = [B \ AB \ \dots \ A^{n-1}B]$$

* Bellman's Principle

- Optimal control problems have an inherently sequential structure.
- Past control inputs affect future states, but future control inputs can not affect past states
- Bellman's Principle "The Principle of Optimality" states the consequence of this for optimal trajectories.



If this path had lower cost starting at X_n , I would have taken it starting from X_0 .

- Sub-trajectories of optimal trajectories must be optimal for appropriately defined sub problem.

* Dynamic Programming

- Bellman's Principle suggests starting from the end of the trajectory and working backwards.
- We've already seen hints of this in the Riccati equation and Pontryagin.
- Define "optimal cost-to-go" a.k.a. "value function" $V(x)$.
- Encodes cost incurred starting from state x at time k if we act optimally

- For LQR:

$$V_N(x) = \frac{1}{2}x^T Q_N x = \frac{1}{2}x^T P_N x$$

- Back up one step and calculate $V_{N-1}(x)$:

$$V_{N-1} = \min_u \underbrace{\frac{1}{2}x_{N-1}^T Q x_{N-1} + \frac{1}{2}u^T R u + V_N(A_{N-1}x_{N-1} + B_{N-1}u)}_{\text{blue underline}}$$

$$= \min_u \frac{1}{2}u^T R u + \frac{1}{2}(Ax_{N-1} + Bu)^T P_N (Ax_{N-1} + Bu)$$

$$\Rightarrow R u + B^T P_N (A x_{N-1} + B u) = 0$$

$$\Rightarrow u_{N-1} = - \underbrace{(R_{N-1} + B_{N-1}^T P_N B_{N-1})^{-1} B_{N-1}^T P_N A_{N-1} x_{N-1}}_{K_{N-1}}$$

- Plug $u = -Kx$ back into *

$$\Rightarrow V_{N-1}(x) = \frac{1}{2}x^T \underbrace{[Q + K^T R K - (A - BK)^T P_N (A - BK)]}_{P_{N-1}} x$$

$$\Rightarrow V_{N-1}(x) = \frac{1}{2}x^T P_{N-1} x$$

- Now we have a backward recursion for K and P that we can iterate until $k=0$

- Just Riccati recursion again!

* Backward Dynamic Programming Algorithms

$$V_N(x) \leftarrow l_N(x)$$

$k \leftarrow N$

while $k > 1$

$$V_{n-1}(x) = \min_{u \in U} [l(x, u) + V_n(f(x, u))]$$

$k \leftarrow k - 1$

end

- If we know $V_n(x)$, the optimal policy is:

$$u_n(x) = \underset{u}{\operatorname{argmax}} [l(x, u) + V_{n+1}(f(x, u))]$$

- DP equations written equivalently in terms of "action-value" or "Q" function:

$$S_n(x, u) = l(x, u) + V_{n+1}(f(x, u))$$

$$\Rightarrow u_n(x) = \underset{u}{\operatorname{argmax}} S_n(x, u)$$

- Usually denoted $Q(x, u)$ but we'll use $S(x, u)$ to avoid confusion with LQR
- Averts need for dynamics model.

* The Curse

- DP is sufficient for a global optimum
- Only tractable for simple problems (LQR or low-dimensional)
- $V(x)$ stays quadratic for LQR, but becomes impossible to write analytically for even simple nonlinear problems.
- Even if we could, $\min_u S(x, u)$ will be non-convex and possibly hard to solve.
- Cost of DP blows up with state dimension due to difficulty of representing $V(x)$

* Then Why DO We Care?

- Approximate DP where $V(x)$ or $S(x, u)$ is represented with a function approximator can work well.
- Forms the basis of modern RL
- DP generalizes to stochastic problems (just wrap everything in expectation operators). Pontryagin does not.

* Finally : What are the Lagrange Multipliers?

- Recall Riccati derivation from QP:

$$\lambda_n = P_n x_n$$

- From DP :

$$V_n(x) = \frac{1}{2} x^T P_n x$$

$$\Rightarrow \lambda_n = D_x V_n(x)$$

- Dynamics multipliers are cost-to-go gradients!

- Carries over to the general non-linear setting (not just LQR).

* Example :

λ_n from QP matches $D_x V_n(x)$ from DP